

# TINY NOTES ON (NOT SO TINY) LANGUAGE MODELS

Leonardo Cotta

Ellison Institute of Technology

*Notes for a two-day mini course on language models  
Ellison Institute of Technology Centre for Doctoral Training  
in the Fundamentals of AI*

# Preamble

“Language models are just next-word (token) predictors”

- What is a language model?

# Preamble

“Language models are just next-word (token) predictors”

- ▶ What is a language model?
- ▶ How does it work?

# Preamble

“Language models are just next-word (token) predictors”

- ▶ What is a language model?
- ▶ How does it work?
- ▶ By the end: you'll be able to build your own

# Preamble

“Language models are just next-word (token) predictors”

- ▶ What is a language model?
- ▶ How does it work?
- ▶ By the end: you'll be able to build your own
- ▶ Where does the magic come from?

# Modeling Language

Let  $x$  be a string and  $p(x)$  the probability of observing it.

- Vocabulary  $V$ : finite set of tokens (e.g., ASCII characters)

# Modeling Language

Let  $x$  be a string and  $p(x)$  the probability of observing it.

- ▶ Vocabulary  $V$ : finite set of tokens (e.g., ASCII characters)
- ▶  $X$ : discrete random variable over  $V^*$

# Modeling Language

Let  $x$  be a string and  $p(x)$  the probability of observing it.

- ▶ Vocabulary  $V$ : finite set of tokens (e.g., ASCII characters)
- ▶  $X$ : discrete random variable over  $V^*$
- ▶ **Goal**: approximate  $p(x)$  given training i.i.d. samples  $x^1, \dots, x^N$

# Autoregressive Models

Chain Rule

By the (Bayes) chain rule:

$$p(x) = \prod_{i=1}^{|x|} p(x_i | x_{<i})$$

# Autoregressive Models

Chain Rule

By the (Bayes) chain rule:

$$p(x) = \prod_{i=1}^{|x|} p(x_i | x_{<i})$$

Learning Objective

Learn parametric model  $p_\theta(x) = \prod_{i=1}^{|x|} p_\theta(x_i | x_{<i})$  by minimizing:

$$\mathcal{L}(\theta) = -\frac{1}{N} \sum_{n=1}^N \log p_\theta(x^n)$$

# Why Autoregressive?

The beauty lies in:

- ▶ Simplifying the joint distribution

# Why Autoregressive?

The beauty lies in:

- ▶ Simplifying the joint distribution
- ▶ Handling both:

# Why Autoregressive?

The beauty lies in:

- ▶ Simplifying the joint distribution
- ▶ Handling both:
  - ▶ Unconditional generation: sample from  $p_\theta(x)$

# Why Autoregressive?

The beauty lies in:

- ▶ Simplifying the joint distribution
- ▶ Handling both:
  - ▶ Unconditional generation: sample from  $p_\theta(x)$
  - ▶ Conditional generation: sample  $y$  from  $p_\theta(y|x)$

# Why Autoregressive?

The beauty lies in:

- ▶ Simplifying the joint distribution
- ▶ Handling both:
  - ▶ Unconditional generation: sample from  $p_\theta(x)$
  - ▶ Conditional generation: sample  $y$  from  $p_\theta(y|x)$
- ▶ Model flexibility + language's flexible data representation

# Why Autoregressive?

The beauty lies in:

- ▶ Simplifying the joint distribution
- ▶ Handling both:
  - ▶ Unconditional generation: sample from  $p_\theta(x)$
  - ▶ Conditional generation: sample  $y$  from  $p_\theta(y|x)$
- ▶ Model flexibility + language's flexible data representation
- ▶ Perfect for: i) scaling and ii) solving general, diverse problems

# Shannon's Model (1948)

Claude Shannon's simple autoregressive model:

- ▶ Count token frequencies after each observed prefix

# Shannon's Model (1948)

Claude Shannon's simple autoregressive model:

- ▶ Count token frequencies after each observed prefix
- ▶ Normalize to get empirical  $\hat{p}(y|x)$

# Shannon's Model (1948)

Claude Shannon's simple autoregressive model:

- ▶ Count token frequencies after each observed prefix
- ▶ Normalize to get empirical  $\hat{p}(y|x)$
- ▶ Compute entropy:  $H(\hat{p}) = -\mathbb{E}_{x \sim \hat{p}}[\log \hat{p}(x)]$

# Shannon's Model (1948)

Claude Shannon's simple autoregressive model:

- ▶ Count token frequencies after each observed prefix
- ▶ Normalize to get empirical  $\hat{p}(y|x)$
- ▶ Compute entropy:  $H(\hat{p}) = -\mathbb{E}_{x \sim \hat{p}}[\log \hat{p}(x)]$

# Shannon's Model (1948)

Claude Shannon's simple autoregressive model:

- ▶ Count token frequencies after each observed prefix
- ▶ Normalize to get empirical  $\hat{p}(y|x)$
- ▶ Compute entropy:  $H(\hat{p}) = -\mathbb{E}_{x \sim \hat{p}}[\log \hat{p}(x)]$

**Strength:** Zero training loss (perfect memorization)

**Weakness:** Zero generalization

# Shannon's Model (1948)

Claude Shannon's simple autoregressive model:

- ▶ Count token frequencies after each observed prefix
- ▶ Normalize to get empirical  $\hat{p}(y|x)$
- ▶ Compute entropy:  $H(\hat{p}) = -\mathbb{E}_{x \sim \hat{p}}[\log \hat{p}(x)]$

**Strength:** Zero training loss (perfect memorization)

**Weakness:** Zero generalization

Compresses *completions* but not *contexts*

# Neural Models

Key Idea

Replace memorization with *inductive learning*

# Neural Models

Key Idea

Replace memorization with *inductive learning*

Learning objective (under infinite data):

$$\theta^* = \arg \min_{\theta} \mathbb{E}_{x \sim p} [-\log p_{\theta}(x)] = \text{KL}(p \| p_{\theta}) + H(p)$$

# Neural Models

Key Idea

Replace memorization with *inductive learning*

Learning objective (under infinite data):

$$\theta^* = \arg \min_{\theta} \mathbb{E}_{x \sim p} [-\log p_{\theta}(x)] = \text{KL}(p \| p_{\theta}) + H(p)$$

Since  $H(p)$  is constant, we minimize  $\text{KL}(p \| p_{\theta})$

# Neural Models

## Key Idea

Replace memorization with *inductive learning*

Learning objective (under infinite data):

$$\theta^* = \arg \min_{\theta} \mathbb{E}_{x \sim p} [-\log p_{\theta}(x)] = \text{KL}(p \| p_{\theta}) + H(p)$$

Since  $H(p)$  is constant, we minimize  $\text{KL}(p \| p_{\theta})$

Neural networks achieve *better* test performance through context compression

# Why Compression Enables Generalization

Consider: "if context ends in token  $a$ , predict token  $b$ "

# Why Compression Enables Generalization

Consider: "if context ends in token  $a$ , predict token  $b$ "

**Shannon:** Store all contexts containing  $a$

# Why Compression Enables Generalization

Consider: "if context ends in token  $a$ , predict token  $b$ "

**Shannon:** Store all contexts containing  $a$

**Neural:** Store the rule itself

# Why Compression Enables Generalization

Consider: "if context ends in token  $a$ , predict token  $b$ "

**Shannon:** Store all contexts containing  $a$

**Neural:** Store the rule itself

**Compression:** Represent data with fewer degrees of freedom by exploiting regularities (symmetries, patterns, structure)

# Why Compression Enables Generalization

Consider: "if context ends in token  $a$ , predict token  $b$ "

**Shannon:** Store all contexts containing  $a$

**Neural:** Store the rule itself

**Compression:** Represent data with fewer degrees of freedom by exploiting regularities (symmetries, patterns, structure)

Successful compression  $\Rightarrow$  extrapolation to unseen contexts

# The Inductive Leap

- With finite data and parameters, the model *must infer* general underlying regularities from observed examples

# The Inductive Leap

- ▶ With finite data and parameters, the model *must infer* general underlying regularities from observed examples
- ▶ **This is induction:** generalizing from particular instances to general principles

# The Inductive Leap

- ▶ With finite data and parameters, the model *must infer* general underlying regularities from observed examples
- ▶ **This is induction:** generalizing from particular instances to general principles

# The Inductive Leap

- ▶ With finite data and parameters, the model *must infer* general underlying regularities from observed examples
- ▶ **This is induction:** generalizing from particular instances to general principles

LLMs are not just interpolating training data—they learn compressive representations that transfer to new contexts

# Pretraining is That Simple

## Remark (Pretraining)

*Pretraining compresses data (e.g., the Internet) into parameters. Chinchilla scaling suggests ~1 parameter per 20 tokens. This creates a “smooth lookup table” that post-training refines for instructions, reasoning, etc.*

# Tokenization

## Motivation

Efficiently represent text as sequences of tokens

# Tokenization

## Motivation

Efficiently represent text as sequences of tokens

## Trade-off:

- ▶ **Characters:** Small vocabulary (128 for ASCII) but long sequences

# Tokenization

## Motivation

Efficiently represent text as sequences of tokens

## Trade-off:

- ▶ **Characters:** Small vocabulary (128 for ASCII) but long sequences
- ▶ **Words:** Shorter sequences but sparse vocabulary ( $|V| \approx 10^5$ )

# Tokenization

## Motivation

Efficiently represent text as sequences of tokens

## Trade-off:

- ▶ **Characters:** Small vocabulary (128 for ASCII) but long sequences
- ▶ **Words:** Shorter sequences but sparse vocabulary ( $|V| \approx 10^5$ )

# Tokenization

## Motivation

Efficiently represent text as sequences of tokens

## Trade-off:

- ▶ **Characters:** Small vocabulary (128 for ASCII) but long sequences
- ▶ **Words:** Shorter sequences but sparse vocabulary ( $|V| \approx 10^5$ )

## Solution: Byte Pair Encoding (BPE)

- ▶ Iteratively merge frequent character pairs

# Tokenization

## Motivation

Efficiently represent text as sequences of tokens

## Trade-off:

- ▶ **Characters:** Small vocabulary (128 for ASCII) but long sequences
- ▶ **Words:** Shorter sequences but sparse vocabulary ( $|V| \approx 10^5$ )

## Solution: Byte Pair Encoding (BPE)

- ▶ Iteratively merge frequent character pairs
- ▶ Guarantees desired vocabulary size

# Tokenization

## Motivation

Efficiently represent text as sequences of tokens

## Trade-off:

- ▶ **Characters:** Small vocabulary (128 for ASCII) but long sequences
- ▶ **Words:** Shorter sequences but sparse vocabulary ( $|V| \approx 10^5$ )

## Solution: Byte Pair Encoding (BPE)

- ▶ Iteratively merge frequent character pairs
- ▶ Guarantees desired vocabulary size
- ▶ Tokens are commonly seen in corpus

# Byte Pair Encoding

```
def bpe(text, num_merges):
    """Byte Pair Encoding tokenizer"""
    vocab = set(text)
    tokens = list(text)

    for _ in range(num_merges):
        # Count adjacent pairs
        pairs = {}
        for i in range(len(tokens)-1):
            pair = (tokens[i], tokens[i+1])
            pairs[pair] = pairs.get(pair, 0) + 1

        if not pairs: break

        # Merge most frequent pair
        best_pair = max(pairs, key=pairs.get)
        new_tokens = []
        i = 0
        while i < len(tokens):
            if i < len(tokens)-1 and \
                (tokens[i], tokens[i+1]) == best_pair:
                new_tokens.append(
                    tokens[i] + tokens[i+1])
                i += 2
            else:
                new_tokens.append(tokens[i])
                i += 1
        tokens = new_tokens
        vocab.add(best_pair[0] + best_pair[1])

    return tokens, vocab
```

# Transformers: Overview

Process sequences up to length  $\ell_{\max}$  with hidden dimension  $d$

## Key Components:

- ▶ Token embeddings

# Transformers: Overview

Process sequences up to length  $\ell_{\max}$  with hidden dimension  $d$

## Key Components:

- ▶ Token embeddings
- ▶ Positional encoding

# Transformers: Overview

Process sequences up to length  $\ell_{\max}$  with hidden dimension  $d$

## Key Components:

- ▶ Token embeddings
- ▶ Positional encoding
- ▶ Multi-head attention

# Transformers: Overview

Process sequences up to length  $\ell_{\max}$  with hidden dimension  $d$

## Key Components:

- ▶ Token embeddings
- ▶ Positional encoding
- ▶ Multi-head attention
- ▶ Feed-forward networks

# Transformers: Overview

Process sequences up to length  $\ell_{\max}$  with hidden dimension  $d$

## Key Components:

- ▶ Token embeddings
- ▶ Positional encoding
- ▶ Multi-head attention
- ▶ Feed-forward networks
- ▶ Residual connections & layer normalization

# Token Embeddings & Positional Encoding

## Token Embeddings

Matrix  $\mathbf{E} \in \mathbb{R}^{|V| \times d}$  where row  $v$  is token  $v$ 's embedding

# Token Embeddings & Positional Encoding

## Token Embeddings

Matrix  $\mathbf{E} \in \mathbb{R}^{|V| \times d}$  where row  $v$  is token  $v$ 's embedding

## Positional Encoding

Captures token order (attention is permutation-equivariant)

Sinusoidal encoding:

$$\text{PE}[i, 2j] = \sin(i/10000^{2j/d})$$

$$\text{PE}[i, 2j + 1] = \cos(i/10000^{2j/d})$$

for position  $i$  and dimension  $j = 0, \dots, d/2 - 1$

# Attention Mechanism

Token gets query  $q \in \mathbb{R}^d$ ; context tokens get keys  $k_t$ , values  $v_t \in \mathbb{R}^d$

# Attention Mechanism

Token gets query  $q \in \mathbb{R}^d$ ; context tokens get keys  $k_t$ , values  $v_t \in \mathbb{R}^d$

$$\alpha_t = \frac{\exp(\langle q, k_t \rangle / \sqrt{d})}{\sum_{t'} \exp(\langle q, k_{t'} \rangle / \sqrt{d})}$$

# Attention Mechanism

Token gets query  $q \in \mathbb{R}^d$ ; context tokens get keys  $k_t$ , values  $v_t \in \mathbb{R}^d$

$$\alpha_t = \frac{\exp(\langle q, k_t \rangle / \sqrt{d})}{\sum_{t'} \exp(\langle q, k_{t'} \rangle / \sqrt{d})}$$

$$\text{Attn}(q, K, V) = \sum_t \alpha_t v_t$$

# Attention Mechanism

Token gets query  $q \in \mathbb{R}^d$ ; context tokens get keys  $k_t$ , values  $v_t \in \mathbb{R}^d$

$$\alpha_t = \frac{\exp(\langle q, k_t \rangle / \sqrt{d})}{\sum_{t'} \exp(\langle q, k_{t'} \rangle / \sqrt{d})}$$

$$\text{Attn}(q, K, V) = \sum_t \alpha_t v_t$$

**Multi-Head Attention:** Run  $h$  attention heads in parallel (each with own  $\mathbf{W}_Q, \mathbf{W}_K, \mathbf{W}_V$ ), concatenate outputs

# Transformer Forward Pass (1/2)

```
def transformer_forward(x, E, L, h, W_Q, W_K,
                      W_V, W_O, W_ff1, W_ff2):
    """
    x: token indices [n]
    E: embedding matrix [V, d]
    L: num layers, h: num heads
    W_Q, W_K, W_V, W_O: attention weights [d, d]
    W_ff1: [d, d_ff], W_ff2: [d_ff, d]
    """
    n, d = len(x), E.shape[1]
    d_k = d // h

    # Embed tokens
    X = E[x, :] # [n, d]

    # Add positional encoding
    PE = zeros(n, d)
    for i in range(n):
        for j in range(d//2):
            PE[i, 2*j] = sin(i / 10000**((2*j)/d))
            PE[i, 2*j+1] = cos(i / 10000**((2*j)/d))
    X = X + PE
```

# Transformer Forward Pass (2/2)

```
# Transformer layers
for layer in range(L):
    Q = X @ W_Q[layer] # [n, d]
    K = X @ W_K[layer]
    V = X @ W_V[layer]

    Q = Q.reshape(n, h, d_k)
    K = K.reshape(n, h, d_k)
    V = V.reshape(n, h, d_k)

    heads = []
    for i in range(h):
        scores = (Q[:, i, :] @ K[:, i, :].T) \
            / sqrt(d_k)
        attn = softmax(scores, axis=1)
        head_out = attn @ V[:, i, :]
        heads.append(head_out)

    H = concat(heads, axis=1)
    X_attn = H @ W_O[layer]
    X = LayerNorm(X + X_attn)

    X_ff = relu(X @ W_ff1[layer]) @ W_ff2[layer]
    X = LayerNorm(X + X_ff)

logits = X @ E.T # [n, V]
return logits
```

# Generation

Given trained model  $p_\theta(x)$ , generate text autoregressively:

- ▶ Sample each token conditioned on prefix

# Generation

Given trained model  $p_\theta(x)$ , generate text autoregressively:

- ▶ Sample each token conditioned on prefix
- ▶ Append it

# Generation

Given trained model  $p_\theta(x)$ , generate text autoregressively:

- ▶ Sample each token conditioned on prefix
- ▶ Append it
- ▶ Repeat

# Sampling Strategies

- **Greedy decoding:**  $x_t = \arg \max_v p_\theta(v|x_{<t})$

# Sampling Strategies

- ▶ **Greedy decoding:**  $x_t = \arg \max_v p_\theta(v|x_{<t})$ 
  - ▶ Deterministic but repetitive

# Sampling Strategies

- ▶ **Greedy decoding:**  $x_t = \arg \max_v p_\theta(v|x_{<t})$ 
  - ▶ Deterministic but repetitive
- ▶ **Temperature sampling:** Scale logits by  $T > 0$

# Sampling Strategies

- ▶ **Greedy decoding:**  $x_t = \arg \max_v p_\theta(v|x_{<t})$ 
  - ▶ Deterministic but repetitive
- ▶ **Temperature sampling:** Scale logits by  $T > 0$ 
  - ▶  $T < 1$  sharpens (more confident)

# Sampling Strategies

- ▶ **Greedy decoding:**  $x_t = \arg \max_v p_\theta(v|x_{<t})$ 
  - ▶ Deterministic but repetitive
- ▶ **Temperature sampling:** Scale logits by  $T > 0$ 
  - ▶  $T < 1$  sharpens (more confident)
  - ▶  $T > 1$  flattens (more random)

# Sampling Strategies

- ▶ **Greedy decoding:**  $x_t = \arg \max_v p_\theta(v|x_{<t})$ 
  - ▶ Deterministic but repetitive
- ▶ **Temperature sampling:** Scale logits by  $T > 0$ 
  - ▶  $T < 1$  sharpens (more confident)
  - ▶  $T > 1$  flattens (more random)
- ▶ **Top-k sampling:** Sample from  $k$  most probable tokens

# Sampling Strategies

- ▶ **Greedy decoding:**  $x_t = \arg \max_v p_\theta(v|x_{<t})$ 
  - ▶ Deterministic but repetitive
- ▶ **Temperature sampling:** Scale logits by  $T > 0$ 
  - ▶  $T < 1$  sharpens (more confident)
  - ▶  $T > 1$  flattens (more random)
- ▶ **Top-k sampling:** Sample from  $k$  most probable tokens
- ▶ **Top-p (nucleus) sampling:** Sample from smallest set whose cumulative probability exceeds  $p$

# Naive Generation

```
def generate_naive(prompt_tokens, max_len, model):
    """
    prompt_tokens: initial token sequence [n]
    max_len: maximum generation length
    model: transformer model
    """
    tokens = prompt_tokens.copy()
    for _ in range(max_len):
        # Run full forward pass on ALL tokens
        logits = model.forward(tokens)

        # Take logits for last position
        next_logits = logits[-1, :]

        # Sample next token
        next_token = sample(next_logits)

        # Append to sequence
        tokens.append(next_token)

        if next_token == EOS_TOKEN:
            break

    return tokens
```

# Naive Generation

```
def generate_naive(prompt_tokens, max_len, model):
    """
    prompt_tokens: initial token sequence [n]
    max_len: maximum generation length
    model: transformer model
    """
    tokens = prompt_tokens.copy()
    for _ in range(max_len):
        # Run full forward pass on ALL tokens
        logits = model.forward(tokens)

        # Take logits for last position
        next_logits = logits[-1, :]

        # Sample next token
        next_token = sample(next_logits)

        # Append to sequence
        tokens.append(next_token)

        if next_token == EOS_TOKEN:
            break

    return tokens
```

**Issue:** At step  $t$ , recompute attention for all positions  $1, \dots, t - 1$ .  
**Cost:**  $O(n^2)$  total

# KV Caching

- **Key insight:** In autoregressive generation, position  $i < t$  only attends to positions  $\leq i$

# KV Caching

- ▶ **Key insight:** In autoregressive generation, position  $i < t$  only attends to positions  $\leq i$
- ▶ Its key and value vectors never change as we generate new tokens

# KV Caching

- ▶ **Key insight:** In autoregressive generation, position  $i < t$  only attends to positions  $\leq i$
- ▶ Its key and value vectors never change as we generate new tokens
- ▶ ⇒ We can cache them!

# Generation with KV Caching (1/2)

```
def generate_with_kv_cache(prompt_tokens,
                           max_len, model):
    """
    prompt_tokens: initial token sequence [n]
    max_len: maximum generation length
    model: transformer with cache support
    """
    tokens = prompt_tokens.copy()

    # Initial forward pass - cache all KV pairs
    logits, kv_cache = \
        model.forward_with_cache(tokens)

    next_token = sample(logits[-1, :])
    tokens.append(next_token)
```

# Generation with KV Caching (2/2)

```
for _ in range(max_len - 1):
    # Forward pass only on NEW token,
    # reuse cached KV pairs
    logits, kv_cache = \
        model.forward_with_cache(
            [next_token],
            kv_cache=kv_cache
        )

    next_token = sample(logits[-1, :])
    tokens.append(next_token)

    if next_token == EOS_TOKEN:
        break

return tokens
```

# KV Caching: Implementation

For each layer and head, maintain cached  $\mathbf{K}, \mathbf{V}$  matrices

When processing token  $t$ :

1. Compute new query  $q_t$ , key  $k_t$ , value  $v_t$

# KV Caching: Implementation

For each layer and head, maintain cached  $\mathbf{K}, \mathbf{V}$  matrices

When processing token  $t$ :

1. Compute new query  $q_t$ , key  $k_t$ , value  $v_t$
2. Append  $k_t$  to cached keys:  $\mathbf{K} \leftarrow [\mathbf{K}; k_t]$

# KV Caching: Implementation

For each layer and head, maintain cached  $\mathbf{K}, \mathbf{V}$  matrices

When processing token  $t$ :

1. Compute new query  $q_t$ , key  $k_t$ , value  $v_t$
2. Append  $k_t$  to cached keys:  $\mathbf{K} \leftarrow [\mathbf{K}; k_t]$
3. Append  $v_t$  to cached values:  $\mathbf{V} \leftarrow [\mathbf{V}; v_t]$

# KV Caching: Implementation

For each layer and head, maintain cached  $\mathbf{K}, \mathbf{V}$  matrices

When processing token  $t$ :

1. Compute new query  $q_t$ , key  $k_t$ , value  $v_t$
2. Append  $k_t$  to cached keys:  $\mathbf{K} \leftarrow [\mathbf{K}; k_t]$
3. Append  $v_t$  to cached values:  $\mathbf{V} \leftarrow [\mathbf{V}; v_t]$
4. Compute attention:  $\text{Attn}(q_t, \mathbf{K}, \mathbf{V})$  using full cached context

# KV Caching: Complexity & Trade-offs

## Complexity

**Complexity:** Reduces per-step cost from  $O(t^2)$  to  $O(t)$

- ▶ Total generation cost:  $O(n^2) \rightarrow O(n)$

# KV Caching: Complexity & Trade-offs

## Complexity

**Complexity:** Reduces per-step cost from  $O(t^2)$  to  $O(t)$

- ▶ Total generation cost:  $O(n^2) \rightarrow O(n)$

## Memory Trade-off

Must store  $\mathbf{K}, \mathbf{V}$  for each layer and head

For model with  $L$  layers,  $h$  heads, dimension  $d$ , sequence length  $n$ :

- ▶ Memory:  $O(L \cdot h \cdot n \cdot d)$

# KV Caching: Complexity & Trade-offs

## Complexity

**Complexity:** Reduces per-step cost from  $O(t^2)$  to  $O(t)$

- ▶ Total generation cost:  $O(n^2) \rightarrow O(n)$

## Memory Trade-off

Must store  $\mathbf{K}, \mathbf{V}$  for each layer and head

For model with  $L$  layers,  $h$  heads, dimension  $d$ , sequence length  $n$ :

- ▶ Memory:  $O(L \cdot h \cdot n \cdot d)$
- ▶ This is why long contexts are expensive!

# Post-Training

A look at the first strategies used in post-training

(In the distant past of 2022)

- ▶ First ideas: **alignment**

# Post-Training

A look at the first strategies used in post-training

(In the distant past of 2022)

- ▶ First ideas: **alignment**
- ▶ Pretrained models just regurgitate compressed Internet text

# Post-Training

A look at the first strategies used in post-training

(In the distant past of 2022)

- ▶ First ideas: **alignment**
- ▶ Pretrained models just regurgitate compressed Internet text
- ▶ Goal: align predictions with human preferences/values

# Post-Training

A look at the first strategies used in post-training

(In the distant past of 2022)

- ▶ First ideas: **alignment**
- ▶ Pretrained models just regurgitate compressed Internet text
- ▶ Goal: align predictions with human preferences/values
- ▶ Make models useful: answer questions, follow instructions, engage meaningfully

# Reinforcement Learning from Human Feedback (RLHF)

## Data

Collect human feedback: ratings, annotations, pairwise preferences

# Reinforcement Learning from Human Feedback (RLHF)

## Data

Collect human feedback: ratings, annotations, pairwise preferences

## Reward Models

Train a model to predict reward scores for generated sequences given prompts

- ▶ Trained on human feedback data

# Reinforcement Learning from Human Feedback (RLHF)

## Data

Collect human feedback: ratings, annotations, pairwise preferences

## Reward Models

Train a model to predict reward scores for generated sequences given prompts

- ▶ Trained on human feedback data
- ▶ Assigns real-valued scores (not probabilities)

# Reinforcement Learning from Human Feedback (RLHF)

## Data

Collect human feedback: ratings, annotations, pairwise preferences

## Reward Models

Train a model to predict reward scores for generated sequences given prompts

- ▶ Trained on human feedback data
- ▶ Assigns real-valued scores (not probabilities)
- ▶ Goal: guide model answers toward higher rewards

# RLHF is Just Bayesian Inference

## The Intuition

Update pretrained model  $p_\theta$  to account for human preferences encoded in reward function  $r$

# RLHF is Just Bayesian Inference

## The Intuition

Update pretrained model  $p_\theta$  to account for human preferences encoded in reward function  $r$

This is Bayesian inference:

- ▶ Prior distribution: pretrained model

# RLHF is Just Bayesian Inference

## The Intuition

Update pretrained model  $p_\theta$  to account for human preferences encoded in reward function  $r$

This is Bayesian inference:

- ▶ Prior distribution: pretrained model
- ▶ Evidence: reward function

# RLHF is Just Bayesian Inference

## The Intuition

Update pretrained model  $p_\theta$  to account for human preferences encoded in reward function  $r$

This is Bayesian inference:

- ▶ Prior distribution: pretrained model
- ▶ Evidence: reward function
- ▶ Compute: posterior distribution

# From Rewards to Distributions

Convert reward function  $r(x)$  into distribution:

$$\pi_{\text{KL-RL}}^*(x) = \frac{1}{Z} \pi_0(x) \exp(r(x)/\beta)$$

# From Rewards to Distributions

Convert reward function  $r(x)$  into distribution:

$$\pi_{\text{KL-RL}}^*(x) = \frac{1}{Z} \pi_0(x) \exp(r(x)/\beta)$$

Where:

- $\pi_0 := p_\theta$  is the prior (pretrained model)

# From Rewards to Distributions

Convert reward function  $r(x)$  into distribution:

$$\pi_{\text{KL-RL}}^*(x) = \frac{1}{Z} \pi_0(x) \exp(r(x)/\beta)$$

Where:

- ▶  $\pi_0 := p_\theta$  is the prior (pretrained model)
- ▶  $\exp(r(x)/\beta)$  is the likelihood (reward evidence, scaled by temperature  $\beta$ )

# From Rewards to Distributions

Convert reward function  $r(x)$  into distribution:

$$\pi_{\text{KL-RL}}^*(x) = \frac{1}{Z} \pi_0(x) \exp(r(x)/\beta)$$

Where:

- ▶  $\pi_0 := p_\theta$  is the prior (pretrained model)
- ▶  $\exp(r(x)/\beta)$  is the likelihood (reward evidence, scaled by temperature  $\beta$ )
- ▶  $Z$  is the partition function (normalization)

# From Rewards to Distributions

Convert reward function  $r(x)$  into distribution:

$$\pi_{\text{KL-RL}}^*(x) = \frac{1}{Z} \pi_0(x) \exp(r(x)/\beta)$$

Where:

- ▶  $\pi_0 := p_\theta$  is the prior (pretrained model)
- ▶  $\exp(r(x)/\beta)$  is the likelihood (reward evidence, scaled by temperature  $\beta$ )
- ▶  $Z$  is the partition function (normalization)
- ▶  $\pi_{\text{KL-RL}}^*$  is the target that balances prior and reward

# The KL-Regularized RL Objective

$$J_{\text{KL-RL}}(\theta) = \mathbb{E}_{x \sim \pi_\theta}[r(x)] - \beta \cdot \text{KL}(\pi_\theta \| \pi_0)$$

# The KL-Regularized RL Objective

$$J_{\text{KL-RL}}(\theta) = \mathbb{E}_{x \sim \pi_\theta}[r(x)] - \beta \cdot \text{KL}(\pi_\theta \| \pi_0)$$

Can be rewritten as minimizing KL divergence to target:

$$J_{\text{KL-RL}}(\theta) \propto -\text{KL}(\pi_\theta \| \pi_{\text{KL-RL}}^*)$$

# The KL-Regularized RL Objective

$$J_{\text{KL-RL}}(\theta) = \mathbb{E}_{x \sim \pi_\theta}[r(x)] - \beta \cdot \text{KL}(\pi_\theta \| \pi_0)$$

Can be rewritten as minimizing KL divergence to target:

$$J_{\text{KL-RL}}(\theta) \propto -\text{KL}(\pi_\theta \| \pi_{\text{KL-RL}}^*)$$

This is *variational inference*: approximating intractable posterior  $\pi_{\text{KL-RL}}^*$  with parametric model  $\pi_\theta$

# Why This View Matters

- KL penalty isn't ad-hoc regularization

# Why This View Matters

- ▶ KL penalty isn't ad-hoc regularization
- ▶ It emerges naturally from casting alignment as posterior inference

# Why This View Matters

- ▶ KL penalty isn't ad-hoc regularization
- ▶ It emerges naturally from casting alignment as posterior inference
- ▶ Prior  $\pi_0$  keeps  $\pi_\theta$  fluent and diverse

# Why This View Matters

- ▶ KL penalty isn't ad-hoc regularization
- ▶ It emerges naturally from casting alignment as posterior inference
- ▶ Prior  $\pi_0$  keeps  $\pi_\theta$  fluent and diverse
- ▶ Reward likelihood guides toward preferred behaviors

# Why This View Matters

- ▶ KL penalty isn't ad-hoc regularization
- ▶ It emerges naturally from casting alignment as posterior inference
- ▶ Prior  $\pi_0$  keeps  $\pi_\theta$  fluent and diverse
- ▶ Reward likelihood guides toward preferred behaviors
- ▶ Temperature  $\beta$  controls trust balance:

# Why This View Matters

- ▶ KL penalty isn't ad-hoc regularization
- ▶ It emerges naturally from casting alignment as posterior inference
- ▶ Prior  $\pi_0$  keeps  $\pi_\theta$  fluent and diverse
- ▶ Reward likelihood guides toward preferred behaviors
- ▶ Temperature  $\beta$  controls trust balance:
  - ▶ High  $\beta$ : stay close to  $\pi_0$

# Why This View Matters

- ▶ KL penalty isn't ad-hoc regularization
- ▶ It emerges naturally from casting alignment as posterior inference
- ▶ Prior  $\pi_0$  keeps  $\pi_\theta$  fluent and diverse
- ▶ Reward likelihood guides toward preferred behaviors
- ▶ Temperature  $\beta$  controls trust balance:
  - ▶ High  $\beta$ : stay close to  $\pi_0$
  - ▶ Low  $\beta$ : follow reward more aggressively

# Tell Us What You Think

<https://forms.office.com/e/J9y561KuDV>

